


# $H_{10}/STO-3G (R, \delta)$ surface: matched surrogate learning curves



Median over 5 random splits (first panel) or 3 model seeds (blocked panels); band is the full range across splits and seeds. All four models share one draw of training indices at every label count.  $\Delta$ -learning wins on interpolation and into the compressed wall, and loses on the stretched block where the correlation energy is the less smooth of the two targets. Interpolation on one  $H_{10}$  surface, not molecular transferability.